

Flink

An overview


Agenda

- What the (**Flink**)?
- **Evolution** of Stream Processing
- **HLA** & Stream processing **Concepts**
- Common **APIs** (DataStream, DataSet, Table API)
- **Run** your own Flink application.
- Best Practices
- **Resources** & Exploring further.

FLINK

DISTRIBUTED DATA
STREAM PROCESSING
FRAMEWORK

Evolution




Traditional DW Architecture

Evolution

All DATA that is created or obtained from a source is essentially a


stream of EVENTS

In order to process a data stream ..




Stateful Streaming Application

Evolution




Analytics Streaming Architecture

HLA




Concepts : DATAFLOW



Concepts : PERFORMANCE

Latency Time taken for an event to be processed.




Throughput A measure of the system's processing capacity.

Concepts: TIME SEMANTICS

Processing Time


vs

Event Time




Concepts: OPERATIONS

- Stateless vs Stateful
- Transformation Operations & Rolling Aggregations
- Window operations (Holistic aggregate, Most recent data ..)



Flink APIs



DataStream

```
val inputStream: DataStream[(Int, Int, Int)] = env.fromElements(  
    (1, 2, 2), (2, 3, 1), (2, 2, 4), (1, 5, 3))  
  
val resultStream: DataStream[(Int, Int, Int)] = inputStream  
    .keyBy(0) // key on first field of the tuple  
    .sum(1)   // sum the second field of the tuple in place
```

DataSet

```
val env = ExecutionEnvironment.getExecutionEnvironment

val text: DataSet[String] = env.fromElements("Who's there?", "I think I hear them.
Stand, ho!")

val counts = text.flatMap { _.toLowerCase.split("\\\\w+") filter { _.nonEmpty } }

.map { (_, 1) }

.groupBy(0)

.sum(1)

counts.print()
```

Table API

```
val tableEnv = ... // create a Table Environment for specific planner batch or streaming

tableEnv.executeSql("CREATE TEMPORARY TABLE table1 ... WITH ( 'connector' = ... )" ) //  
create an input Table

tableEnv.executeSql("CREATE TEMPORARY TABLE outputTable ... WITH ( 'connector' = ... )" ) //  
register an output Table

val table2 = tableEnv.from("table1").select(...) // create a Table from a Table API query

val table3 = tableEnv.sqlQuery("SELECT ... FROM table1 ..." ) // create a Table from a SQL  
query

// emit a Table API result Table to a Table Sink, same for SQL result

val tableResult = table2.executeInsert("outputTable")
```

FLINK RUN

Best Practices

- Build a solid CICD pipeline on your infra.
- Leverage event streaming platforms such as Kafka, Kinesis
- Periodic Checkpoints & Savepoints.
- Choose wisely between Event Time and Process Time.
- Develop your own configurable app framework.
- Provision for batch using the same pipeline.
- Setup Monitoring and alert.

References

- <https://ci.apache.org/projects/flink/flink-docs-release-1.12/>
- <https://www.flink-forward.org/>
- **Stream Processing with Apache Flink** - by Fabian Hueske; Vasiliki Kalavri
- <https://eng.uber.com/kappa-architecture-data-stream-processing/>
- Flink Courses (LinkedIn Learning)
- <https://github.com/nupsea/flink-wlm>

?'s

Thank you

- eanups@yahoo.com